


Candidate Result Figures



Prompt compression and local LLM traffic

Preliminary baseline experiment


F1 — Compression reaches distinct token budgets




F2 — Upload traffic falls with prompt compression




F3 — Total traffic reduction depends on the workload




Event summary




F4 — QA traffic is dominated by prompt upload




F5 — Generated summaries dominate bidirectional traffic





F6 — TLS and HTTP/3 transfer similar total bytes





F7 — HTTP/3 uses more packets for short QA exchanges





F8 — Compression shortens QA latency more strongly




F9 — Input tokens predict client-to-server traffic





F10 — Paired samples expose workload-dependent gains



F11 — Summary output limits end-to-end traffic savings



F12 — Absolute traffic separates the two workloads



F13 — Latency reduction follows the input-dominated path

